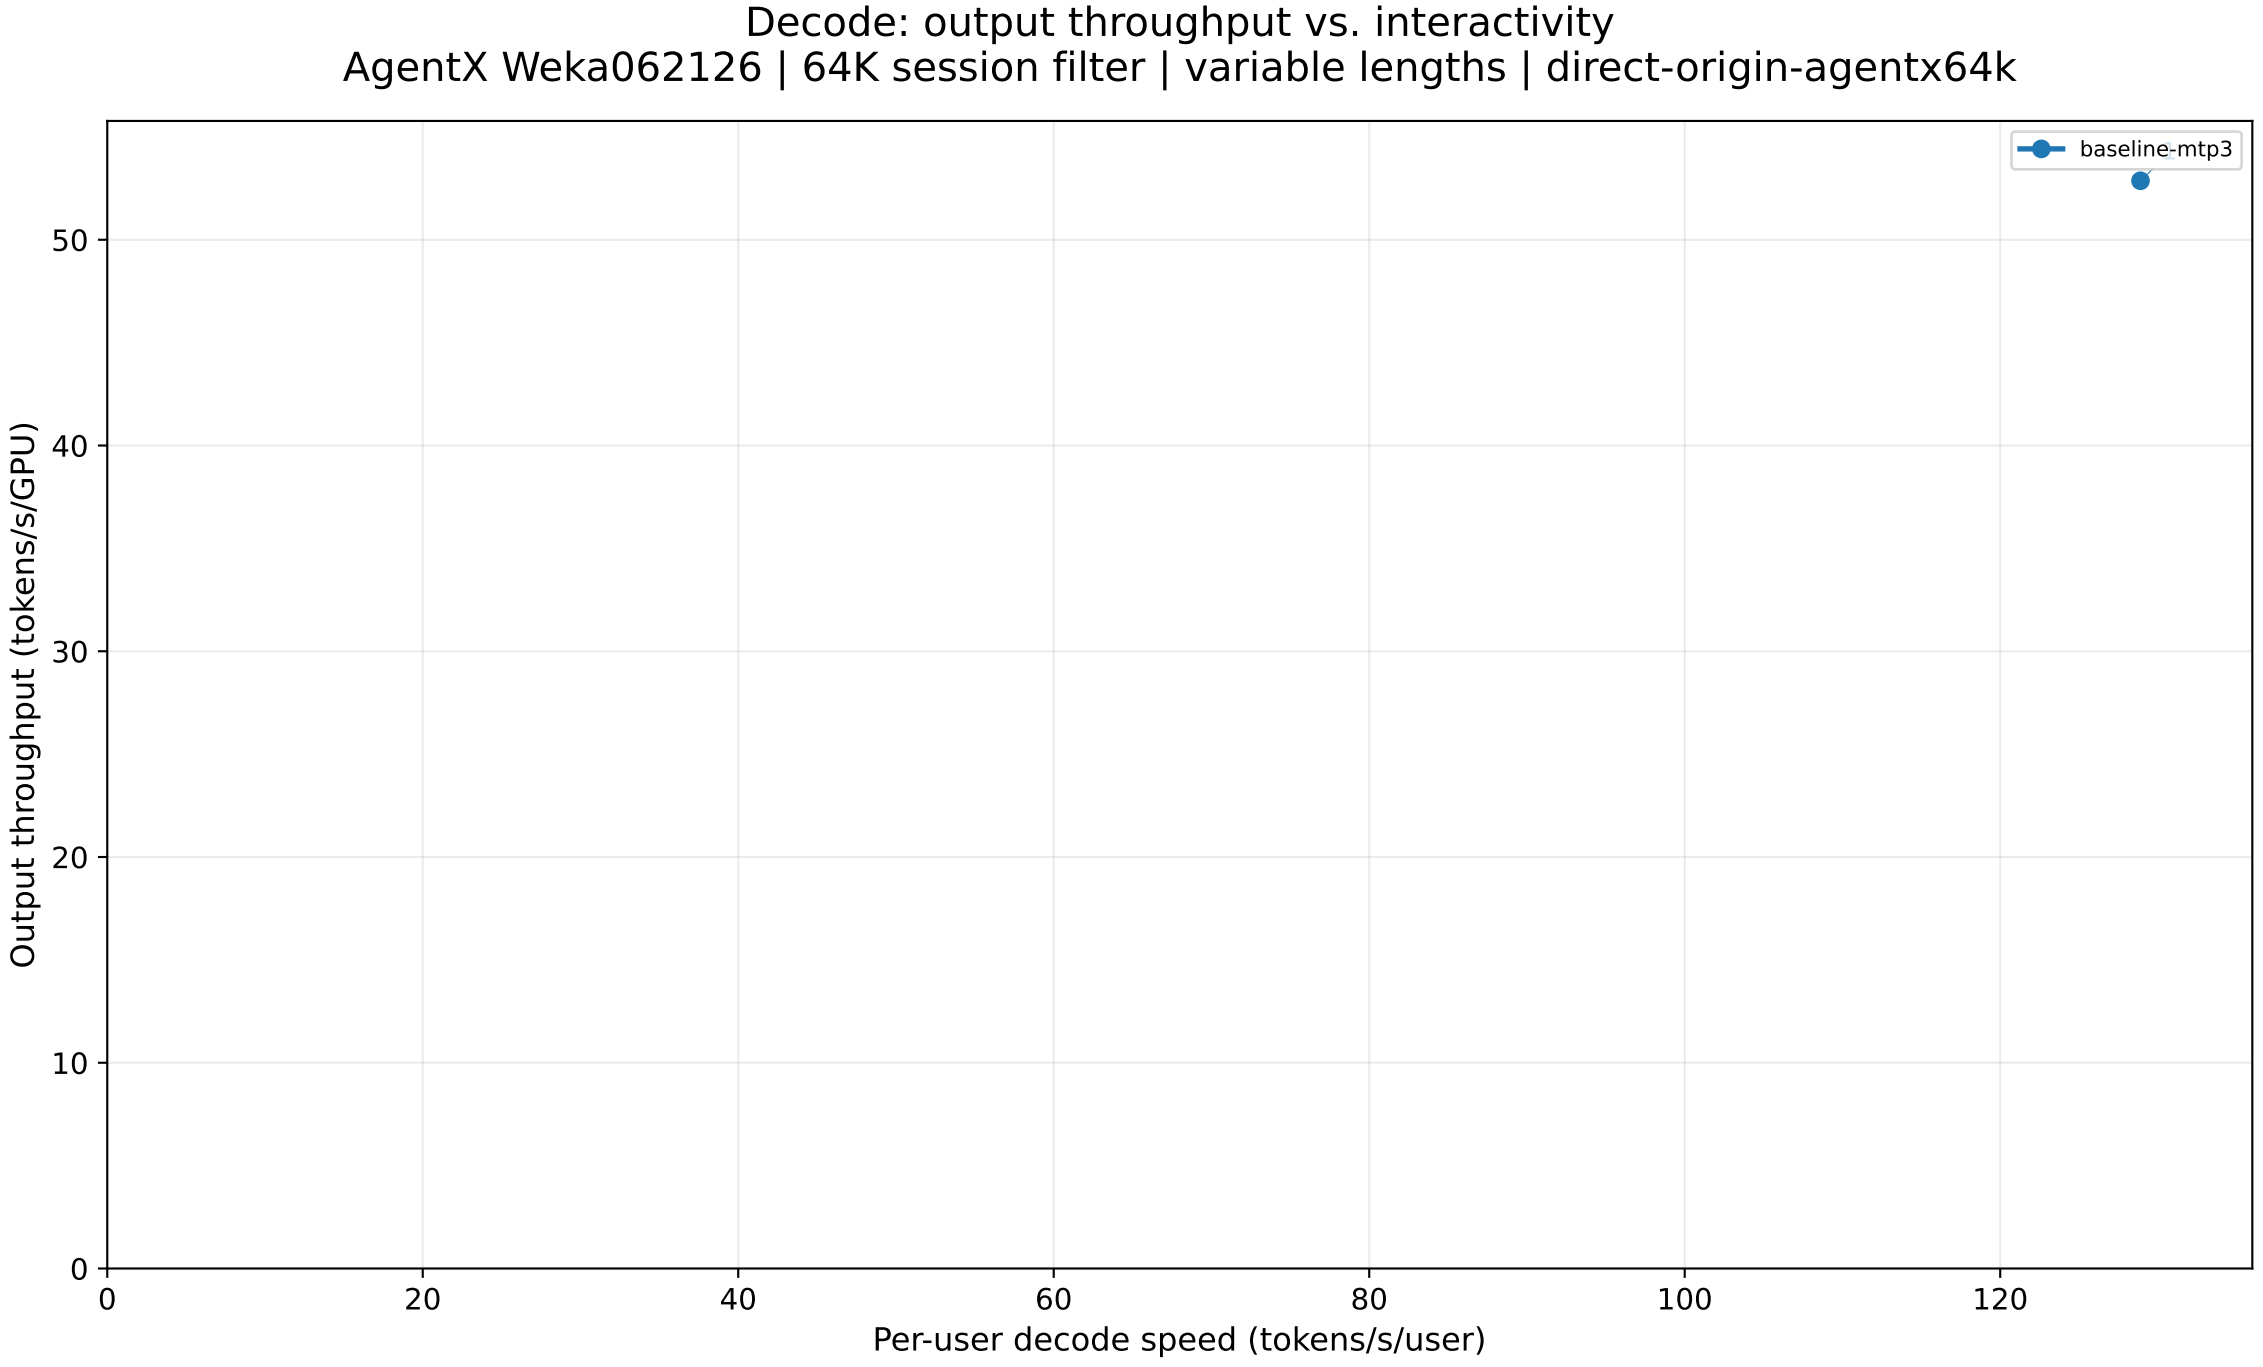
Decode: output throughput vs. interactivity
AgentX Weka062126 | 64K session filter | variable lengths | direct-origin-agentx64k
baseline-mtp3
Output throughput (tokens/s/GPU)
0
10
20
30
40
50
0
20
40
60
80
100
120
Per-user decode speed (tokens/s/user)
Qwen3.8-27B-NVFP4 | 1 RTX PRO 6000 Server Edition | labels: requested client concurrency
AgentX x: reported per-user decode rate. y: raw request-span output throughput.
Progress view: measured points only; missing concurrencies remain unmeasured. Whiskers: observed repeat range, not confidence intervals.