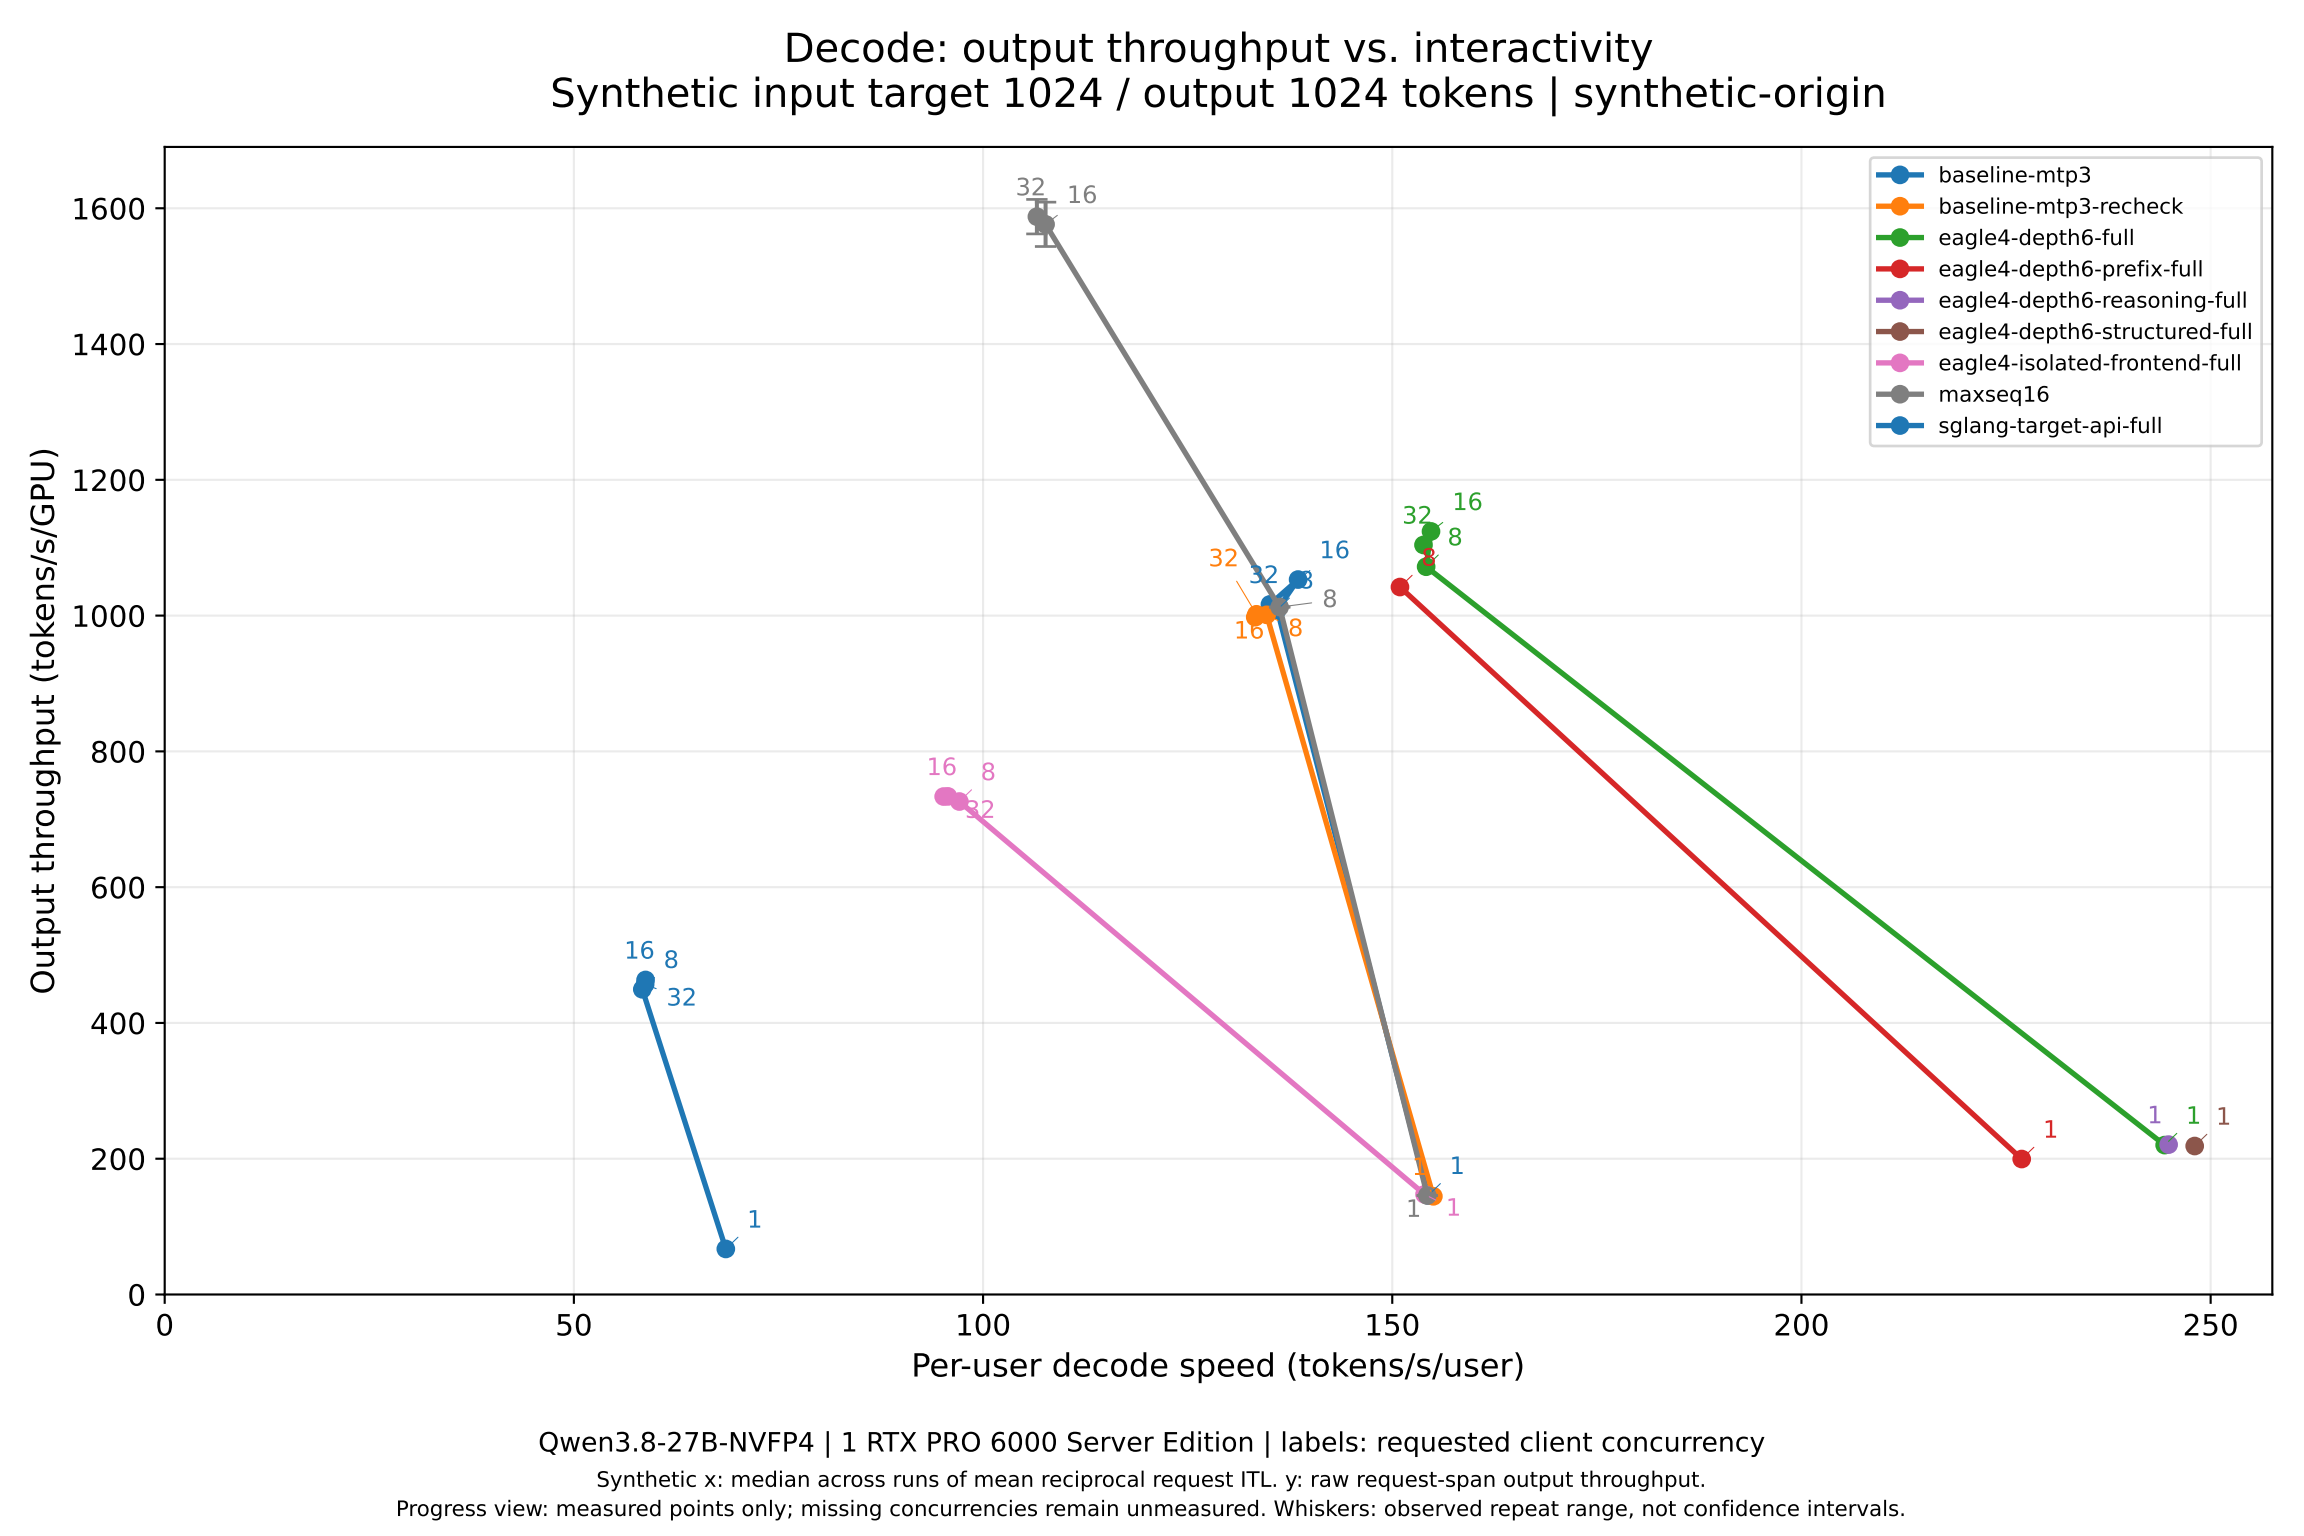
Decode: output throughput vs. interactivity
Synthetic input target 1024 / output 1024 tokens | synthetic-origin
Output throughput (tokens/s/GPU)
Per-user decode speed (tokens/s/user)
baseline-mtp3
baseline-mtp3-recheck
eagle4-depth6-full
eagle4-depth6-prefix-full
eagle4-depth6-reasoning-full
eagle4-depth6-structured-full
eagle4-isolated-frontend-full
maxseq16
sglang-target-api-full
Qwen3.8-27B-NVFP4 | 1 RTX PRO 6000 Server Edition | labels: requested client concurrency
Synthetic x: median across runs of mean reciprocal request ITL. y: raw request-span output throughput.
Progress view: measured points only; missing concurrencies remain unmeasured. Whiskers: observed repeat range, not confidence intervals.